

# Safari | O'Reilly

## CLASS PREPARATION

### Building Distributed Pipelines for Data Science Using Kafka, Spark, and Cassandra

#### Table of contents

<b>Environment</b>	<b>1</b>
<b>Prerequisites</b>	<b>1</b>
Import the Virtual Machine	3
If you are stuck in Virtual Machine	6
<b>Connecting to the Virtual Machine</b>	<b>7</b>
<b>Connect to Spark notebook</b>	<b>8</b>
<b>Stop the Virtual Machine</b>	<b>8</b>

## Environment

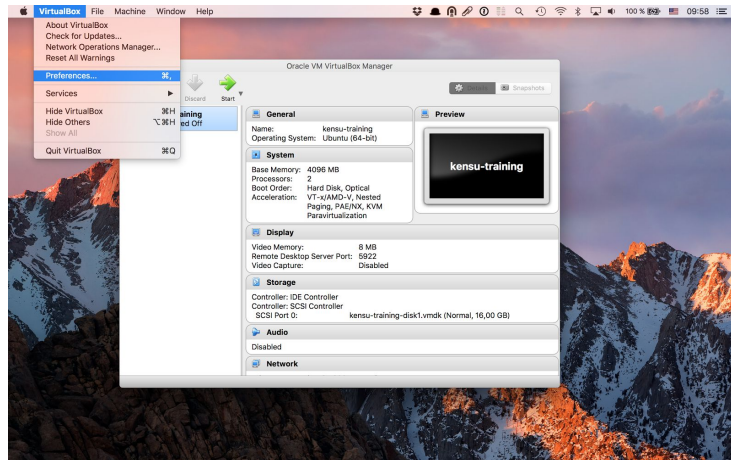
For the training, we'll be using as the simplest environment to run most of the pipeline. This environment will be available from a single virtual machine image with Virtualbox.

## Prerequisites


To use the training Virtual Machine, install [Virtual Box](https://www.virtualbox.org) ( <https://www.virtualbox.org> ) to your host.

Once it's installed, you need to make some customization before importing the training image. You first need to configure a "Host Only Network" to ensure you're able to communicate with the VM from your Host. If you have already one, ensure the DHCP server is enabled.

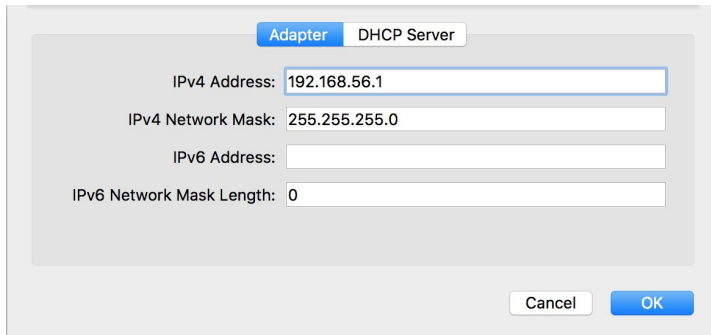
To do this, open VirtualBox Preferences:



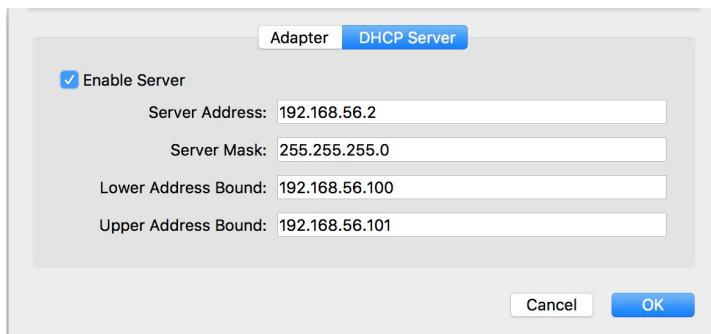
Click on the “Network” tab, then click on “Host-Only Networks” and finally click on the “+” button on the right.



The “vboxnet0” will be created. You also need to enable the DHCP server by selecting the “vboxnet0” and clicking on the screwdriver button.



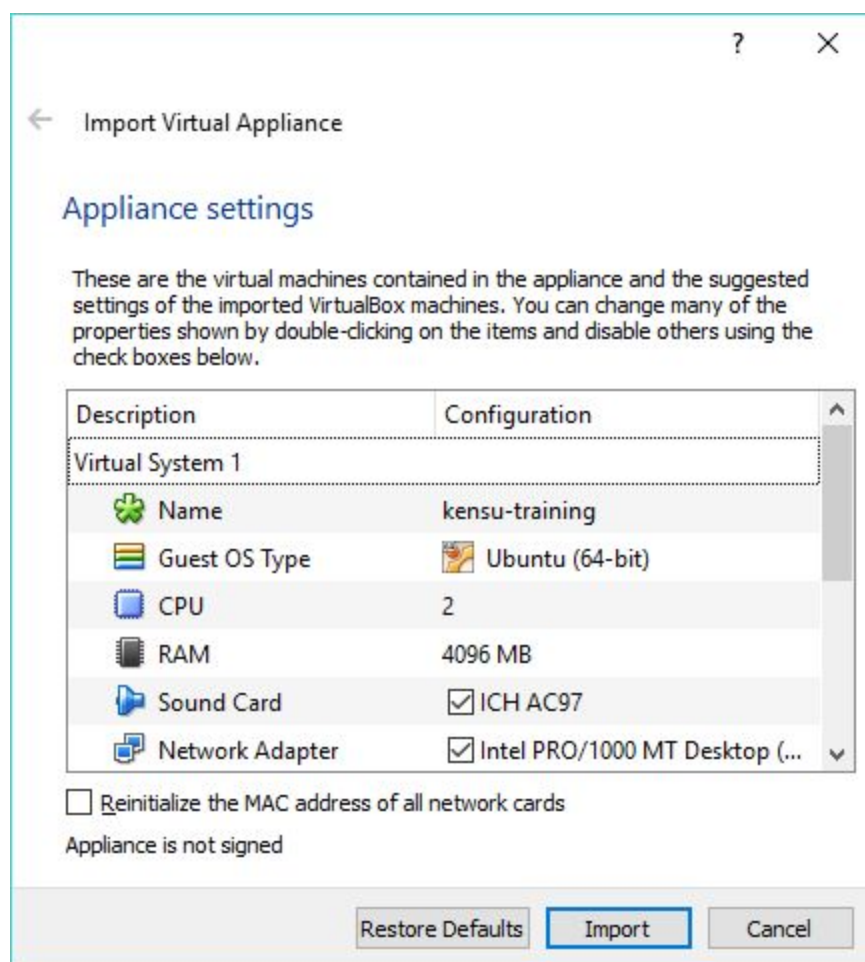
Select the “DCHP Server” tab, ensure the checkbox “Enable Server” is ticked.




## Import the Virtual Machine

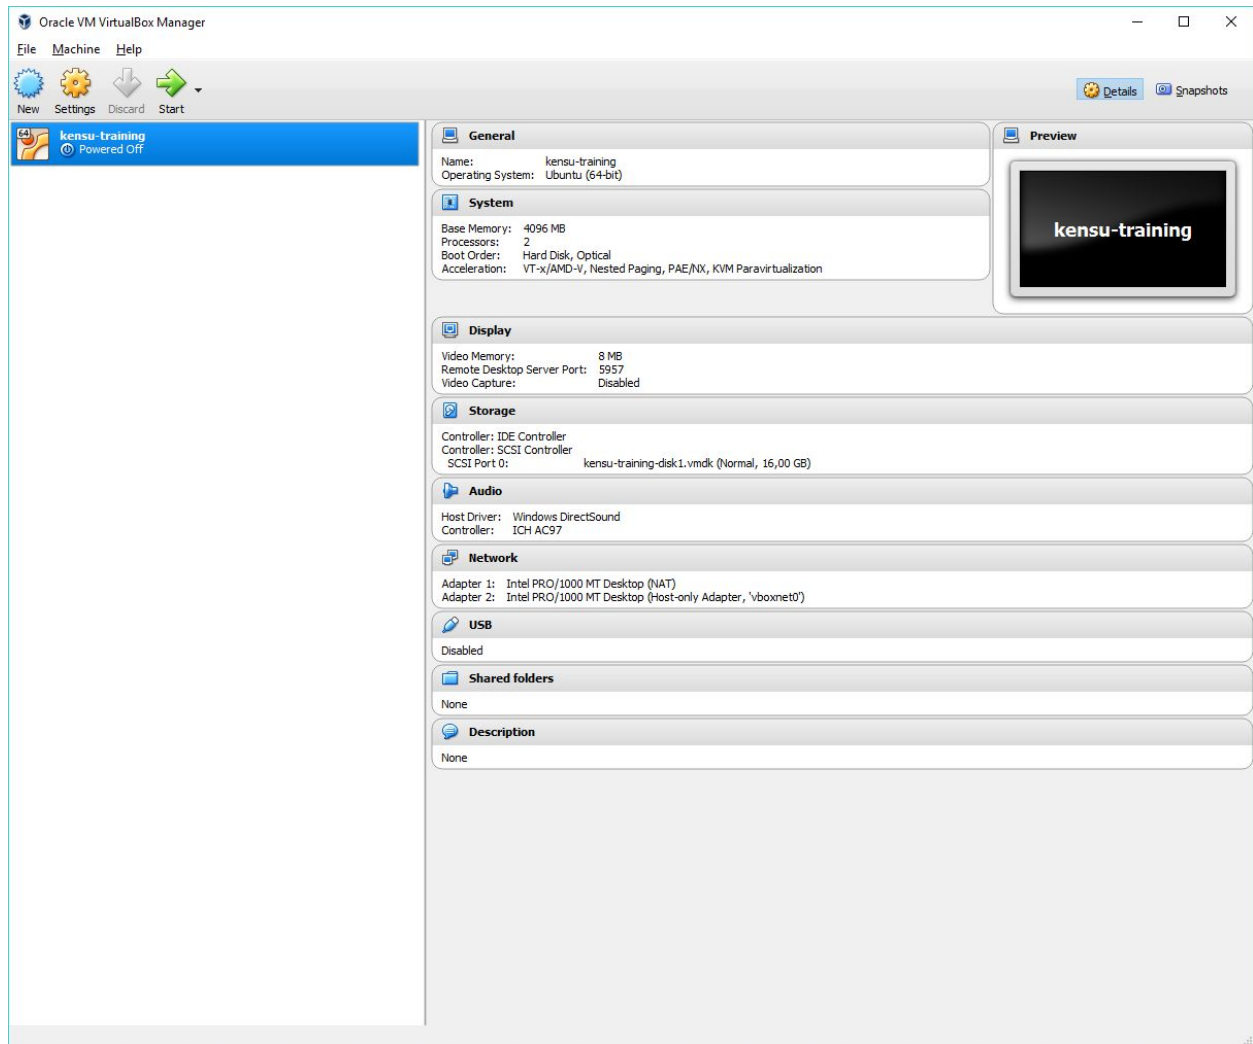
If vbox is running, shut down it.

Download OVA file from <http://kensuio-training.s3.amazonaws.com/kensu-training-00.ova> then double-click on it.




Click on import.



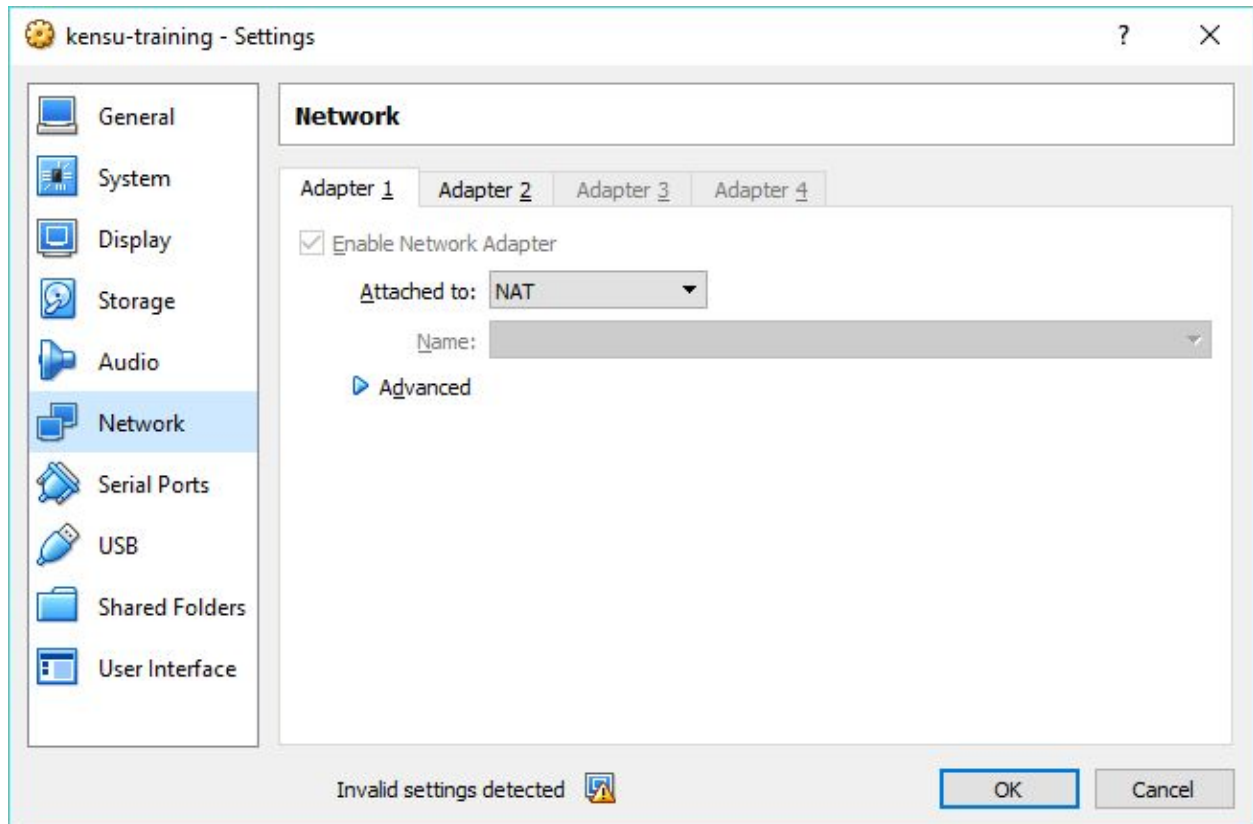


Select the VM and the click to “Start”

If you’re are running Windows, you’ll probably have this error screen:



Click on “Change Network Settings”




Then, click on “OK” button, no changes are required.

The virtual machine should be now running:

Now you can login with kensu as username, and kensu as password.

## If you are stuck in Virtual Machine

**REMEMBER THE DEFAULT KEY TO ESCAPE FROM THE VIRTUAL MACHINE IS **RIGHT CONTROL****



If you're stuck in the console, type “sudo halt -p”

## Connecting to the Virtual Machine

After starting the Training VM in the step above, a shell will open, allowing you to connect to the VM.

Use the credentials:

User: kensu

Password: kensu

You will see a screen like this:

```
kensu-training [Running] - Oracle VM VirtualBox
Ubuntu 16.10 ubuntu tty1
Please connect with login: kensu, password: kensu
You can run "update-notebooks.sh" to update notebook from the latest version.
ubuntu login: kensu
Password:
Welcome to Ubuntu 16.10 (GNU/Linux 4.8.0-22-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage

You can now connect to Spark-notebook from your browser with: http://192.168.56.101:9001
kensu@ubuntu:~$ _
```

It is also recommended that you run “update-notebooks.sh” to get the latest updates to the notebooks used during the session.

## Connect to Spark notebook

Using the output of the shell like illustrated above, copy the URL provided to access the notebook from your browser. (Chrome is strongly recommended, Firefox is OK. IE is not supported)

It's not absolutely necessary to login into the VM. It will often use the default IP address: [192.168.56.101](http://192.168.56.101), so after pressing “start” and waiting for the VM to boot, you could directly access the notebook at the URL: <http://192.168.56.101:9001/>

Note that this default IP Address might change. We strongly recommend login into the shell to get the exact address.

From the VM, we can always inspect the assigned local IP address issuing the following command:

```
kensu@ubuntu:~$ ip addr show dev eth1
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
state UP group default qlen 1000
    link/ether 08:00:27:a8:2e:32 brd ff:ff:ff:ff:ff:ff
    inet 192.168.56.101/24 brd 192.168.56.255 scope global eth1
        valid_lft forever preferred_lft forever
    inet6 fe80::a00:27ff:fea8:2e32/64 scope link
        valid_lft forever preferred_lft forever
```

In this example, the IP address is the text in bold.

Then, try to connect from your browser to <http://192.168.56.101:9001/> (replace the IP address with the one you identified)

## Stop the Virtual Machine

Connect to the Virtual Machine with login kensu, password kensu, and execute the following command:

```
kensu@ubuntu:~$: sudo halt -p
```