influx|days

InfluxDB 101

Michael Desa / Software Engineer

Agenda

By the end of this session, users should be able to...

- Define time-series data
- Describe what InfluxDB is and its relationship to InfluxData
- Explain the InfluxDB data model
- Reason about the impact of the schema in an instance

What is time-series data?


Dow Jones Industrial Average (^DJI)
DJI - DJI Real Time Price. Currency in USD
Add to watchlist
24,834.96 +33.60 (+0.14%)
As of 2:56PM EST. Market open.
Summary
Chart
Options
Components
Historical Data
1D 5D 1M 6M YTD 1Y 5Y Max
Full screen
28,000.00
24,834.81
22,000.00
19,000.00
Mar 8, 17
Sep 7, 17


Alert Processing
99% Alert Processing Latency
80.5 ms
99% Actor Recovery Latency
99.0 ms
99% Service Latency
Chat Connects & Disconnects
Cassandra Reads & Timeouts
Incident & Notification Creation
99% Chat Request Latency
143.4 ms
99% Delivery Latency
9.90 ms
Top 10 Open Incident Count


Disk Ops Write
1.5 K
1.0 K
500
0
Disk Used
233 GB
186 GB
140 GB
93 GB
47 GB
0 B
Short Term Load
20
15
10
5
0
Influxd RSS
60 GB
50 GB
40 GB
30 GB
20 GB
10 GB
0

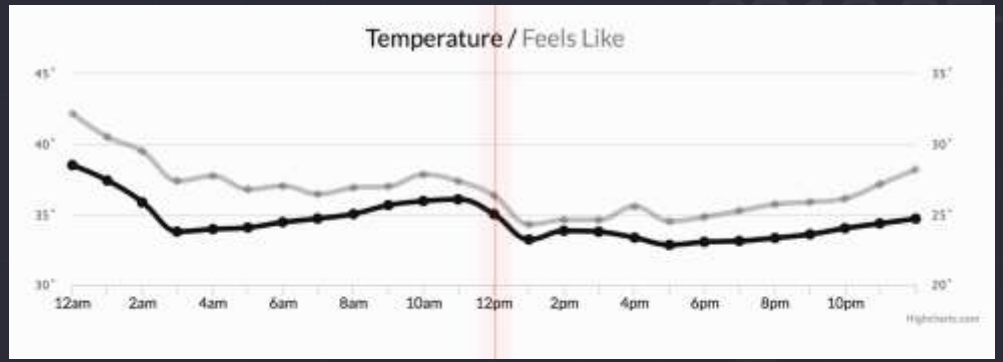
Temperature / Feels Like
45°
40°
35°
30°
35°
30°
25°
20°
12am
2am
4am
6am
8am
10am
12pm
2pm
4pm
6pm
8pm
10pm

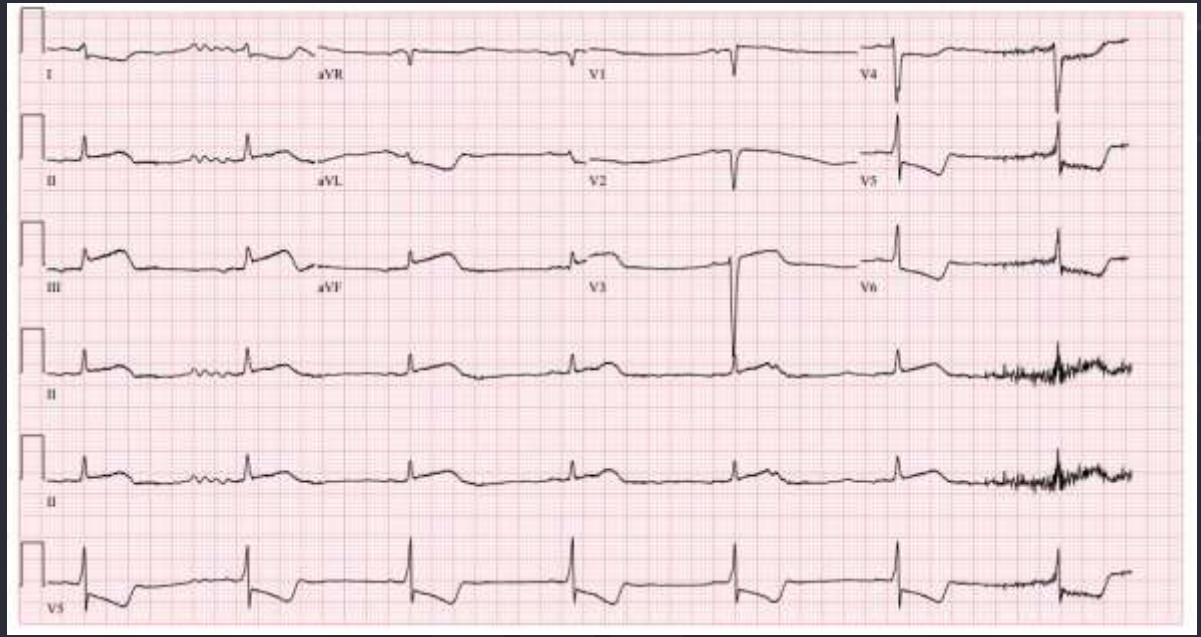
I
aVR
V1
V4
II
aVL
V2
V5
III
aVF
V3
V6
II
II
V5

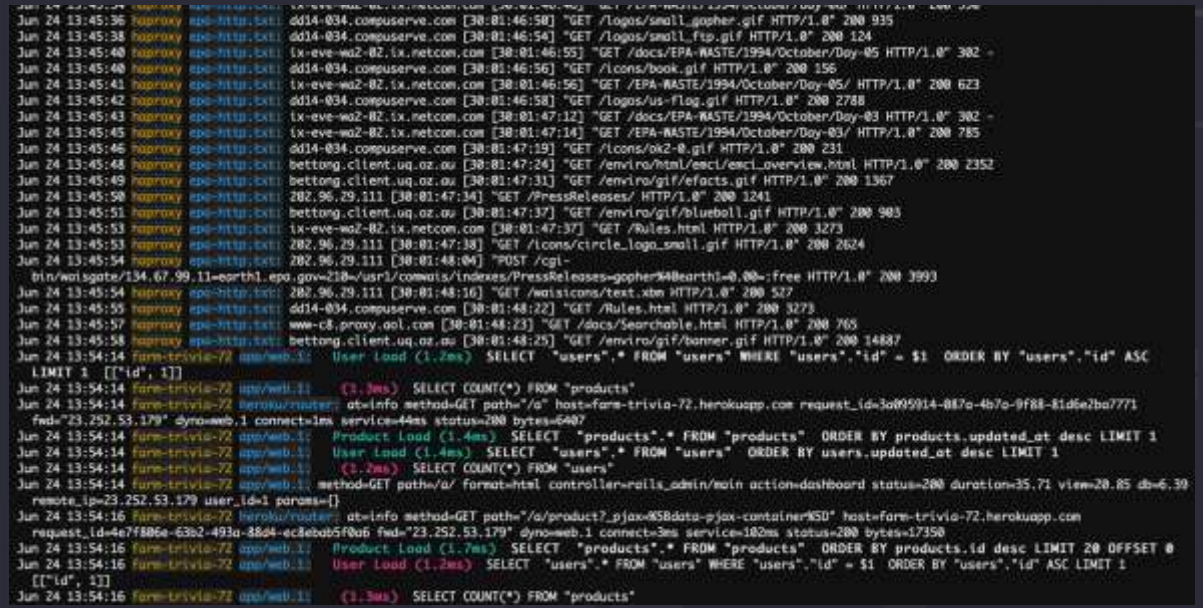
Jun 24 13:45:36 haproxy epo-http.txt: dd14-034.compuserve.com [30:01:46:50] "GET /logos/small_gopher.gif HTTP/1.0" 200 935
Jun 24 13:45:38 haproxy epo-http.txt: dd14-034.compuserve.com [30:01:46:54] "GET /logos/small_ftp.gif HTTP/1.0" 200 124
Jun 24 13:45:40 haproxy epo-http.txt: ix-eve-wa2-02.ix.netcom.com [30:01:46:55] "GET /docs/EPA-WASTE/1994/October/Day-05 HTTP/1.0" 302 -
Jun 24 13:45:40 haproxy epo-http.txt: dd14-034.compuserve.com [30:01:46:56] "GET /icons/book.gif HTTP/1.0" 200 156
Jun 24 13:45:41 haproxy epo-http.txt: ix-eve-wa2-02.ix.netcom.com [30:01:46:56] "GET /EPA-WASTE/1994/October/Day-05/ HTTP/1.0" 200 623
Jun 24 13:45:42 haproxy epo-http.txt: dd14-034.compuserve.com [30:01:46:58] "GET /logos/us-flag.gif HTTP/1.0" 200 2788
Jun 24 13:45:43 haproxy epo-http.txt: ix-eve-wa2-02.ix.netcom.com [30:01:47:12] "GET /docs/EPA-WASTE/1994/October/Day-03 HTTP/1.0" 302 -
Jun 24 13:45:45 haproxy epo-http.txt: ix-eve-wa2-02.ix.netcom.com [30:01:47:14] "GET /EPA-WASTE/1994/October/Day-03/ HTTP/1.0" 200 785
Jun 24 13:45:46 haproxy epo-http.txt: dd14-034.compuserve.com [30:01:47:19] "GET /icons/ok2-0.gif HTTP/1.0" 200 231
Jun 24 13:45:48 haproxy epo-http.txt: bettong.client.uq.oz.au [30:01:47:24] "GET /enviro/html/emci/emci_overview.html HTTP/1.0" 200 2352
Jun 24 13:45:49 haproxy epo-http.txt: bettong.client.uq.oz.au [30:01:47:31] "GET /enviro/gif/efacts.gif HTTP/1.0" 200 1367
Jun 24 13:45:50 haproxy epo-http.txt: 202.96.29.111 [30:01:47:34] "GET /PressReleases/ HTTP/1.0" 200 1241
Jun 24 13:45:51 haproxy epo-http.txt: bettong.client.uq.oz.au [30:01:47:37] "GET /enviro/gif/blueball.gif HTTP/1.0" 200 903
Jun 24 13:45:53 haproxy epo-http.txt: ix-eve-wa2-02.ix.netcom.com [30:01:47:37] "GET /Rules.html HTTP/1.0" 200 3273
Jun 24 13:45:53 haproxy epo-http.txt: 202.96.29.111 [30:01:47:38] "GET /icons/circle_logo_small.gif HTTP/1.0" 200 2624
Jun 24 13:45:54 haproxy epo-http.txt: 202.96.29.111 [30:01:48:04] "POST /cgi-
bin/woisgate/134.67.99.11=earth1.epa.gov=210=/usr1/conmais/indexes/PressReleases=gopher%40earth1=0.00=:free HTTP/1.0" 200 3993
Jun 24 13:45:54 haproxy epo-http.txt: 202.96.29.111 [30:01:48:16] "GET /waisicons/text.xbm HTTP/1.0" 200 527
Jun 24 13:45:55 haproxy epo-http.txt: dd14-034.compuserve.com [30:01:48:22] "GET /Rules.html HTTP/1.0" 200 3273
Jun 24 13:45:57 haproxy epo-http.txt: www-c8.proxy.aol.com [30:01:48:23] "GET /docs/Searchable.html HTTP/1.0" 200 765
Jun 24 13:45:58 haproxy epo-http.txt: bettong.client.uq.oz.au [30:01:48:25] "GET /enviro/gif/banner.gif HTTP/1.0" 200 14887
Jun 24 13:54:14 form-trivia-72 app/web.1: User Load (1.2ms) SELECT "users".* FROM "users" WHERE "users"."id" = $1 ORDER BY "users"."id" ASC
LIMIT 1 [["id", 1]]
Jun 24 13:54:14 form-trivia-72 app/web.1: (1.3ms) SELECT COUNT(*) FROM "products"
Jun 24 13:54:14 form-trivia-72 heroku/router: at=info method=GET path="/a" host=form-trivia-72.herokuapp.com request_id=3a095014-887a-4b7a-9f88-81d6e2ba7771
fwd="23.252.53.179" dyno=web.1 connect=1ms service=44ms status=200 bytes=6407
Jun 24 13:54:14 form-trivia-72 app/web.1: Product Load (1.4ms) SELECT "products".* FROM "products" ORDER BY products.updated_at desc LIMIT 1
Jun 24 13:54:14 form-trivia-72 app/web.1: User Load (1.4ms) SELECT "users".* FROM "users" ORDER BY users.updated_at desc LIMIT 1
Jun 24 13:54:14 form-trivia-72 app/web.1: (1.2ms) SELECT COUNT(*) FROM "users"
Jun 24 13:54:14 form-trivia-72 app/web.1: method=GET path=/a/ format=html controller=rails_admin/main action=dashboard status=200 duration=35.71 view=20.85 db=6.39
remote_ip=23.252.53.179 user_id=1 params={}
Jun 24 13:54:16 form-trivia-72 heroku/router: at=info method=GET path="/a/product?_pjax=%5Bdata-pjax-container%5D" host=form-trivia-72.herokuapp.com
request_id=4e7f8f6e-63b2-493a-8fd4-ec8eab5f0a6 fwd="23.252.53.179" dyno=web.1 connect=3ms service=102ms status=200 bytes=17350
Jun 24 13:54:16 form-trivia-72 app/web.1: Product Load (1.7ms) SELECT "products".* FROM "products" ORDER BY products.id desc LIMIT 20 OFFSET 0
Jun 24 13:54:16 form-trivia-72 app/web.1: User Load (1.2ms) SELECT "users".* FROM "users" WHERE "users"."id" = $1 ORDER BY "users"."id" ASC LIMIT 1
[["id", 1]]
Jun 24 13:54:16 form-trivia-72 app/web.1: (1.3ms) SELECT COUNT(*) FROM "products"

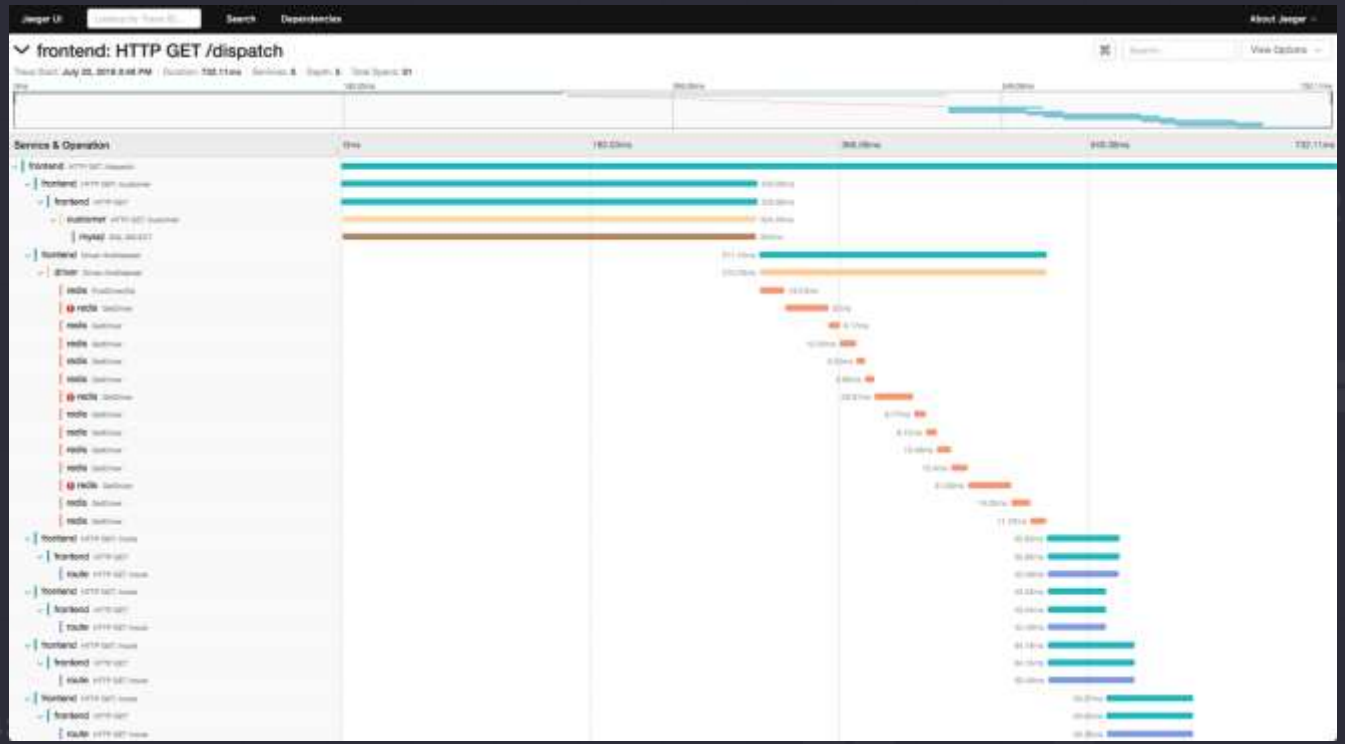
Jaeger UI
Search
Dependencies
About Jaeger
frontend: HTTP GET /dispatch
Service & Operation

Regular vs Irregular Time-Series

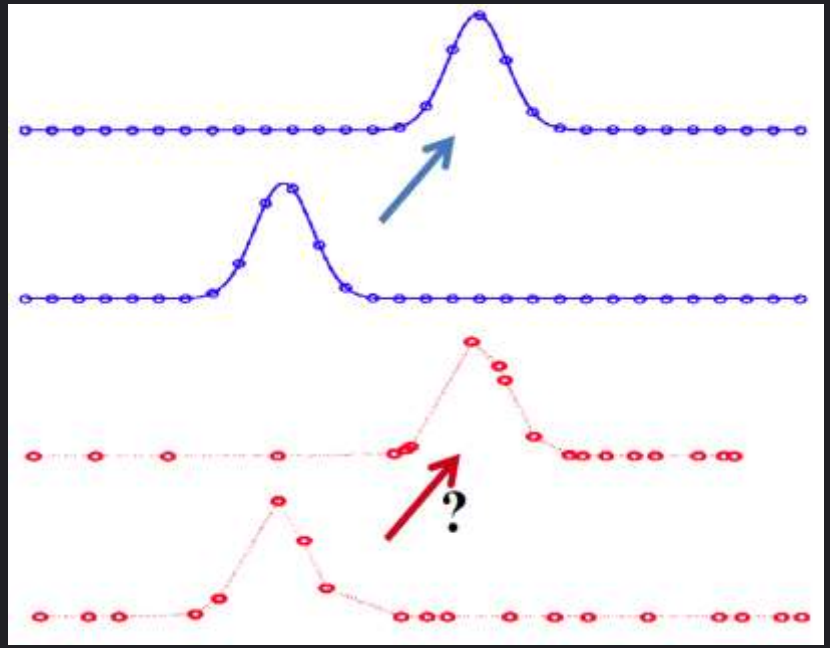

Metrics vs Events

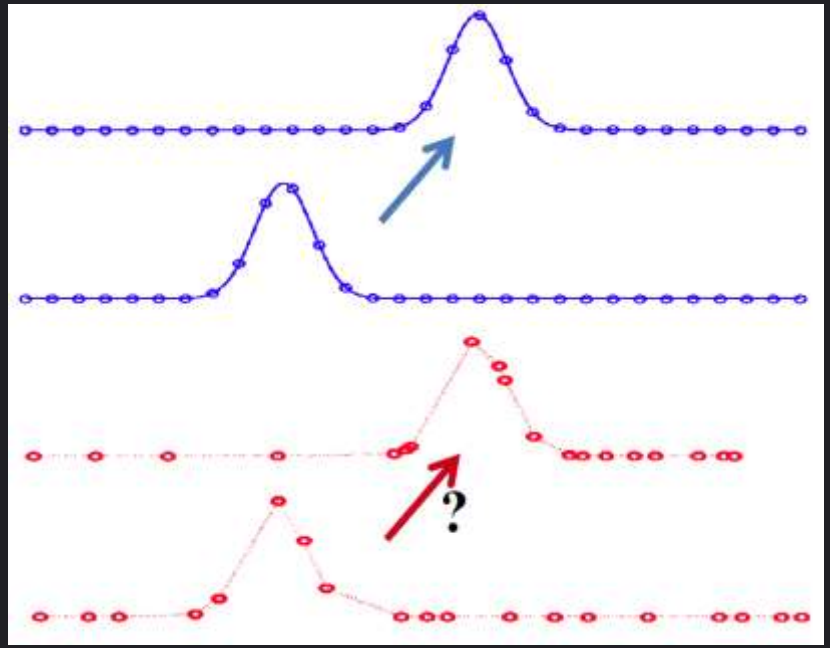

Question: Time-Series?

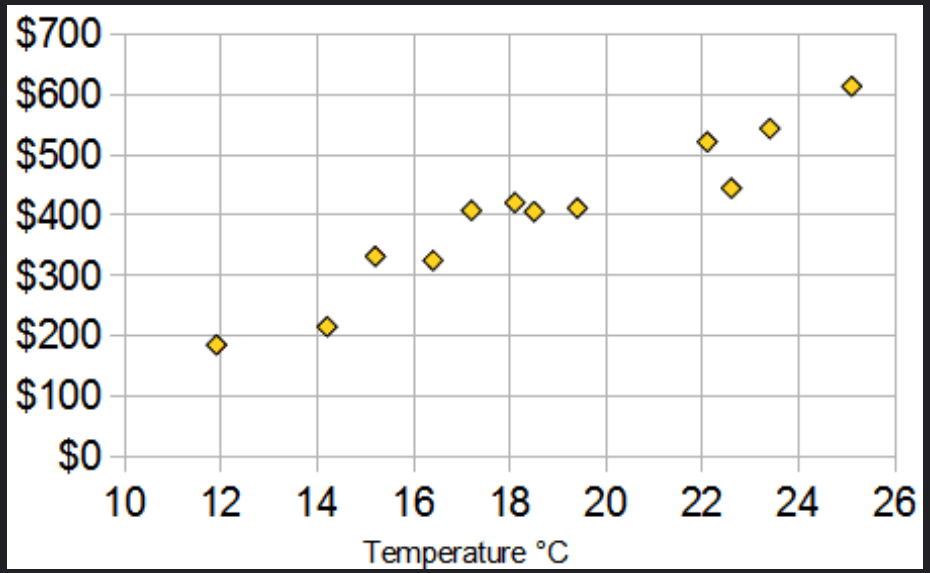

Question: Time-Series?

What is a time-series database (tsdb)?

Time-series Database

- A database where you manage and store time-series data
- Efficiently handles time-series data
- Supports time based queries

Rank			DBMS	Database Model	Score		
May 2019	Apr 2019	May 2018			May 2019	Apr 2019	May 2018
1.	1.	1.	InfluxDB	Time Series	18.08	+0.86	+7.08
2.	2.	2.	Kdb+	Time Series, Multi-model	5.60	-0.25	+2.52
3.	3.	↑ 4.	Graphite	Time Series	3.23	+0.10	+0.96
4.	4.	↑ 6.	Prometheus	Time Series	3.11	+0.20	+1.99
5.	5.	↓ 3.	RRDtool	Time Series	2.90	+0.19	+0.21
6.	6.	↓ 5.	OpenTSDB	Time Series	2.47	+0.10	+0.85
7.	7.	7.	Druid	Multi-model	1.69	+0.04	+0.67
8.	8.	↑ 18.	TimescaleDB	Time Series, Multi-model	1.16	+0.21	+1.12
9.	9.	↓ 8.	KairosDB	Time Series	0.54	-0.09	+0.12
10.	10.	↓ 9.	eXtremeDB	Multi-model	0.38	-0.02	+0.07

Why couldn’t I just use [insert db]?

But, time-series is not just a database problem

Time-series problems

- Visualizing your data
- Alerting you data
- Processing your data
- Taking action based on your data

What is InfluxDB/InfluxData?

InfluxData History

- 2012 - Errplane
- 2014 - InfluxDB is born
- 2015 - Transition to InfluxData
  - A platform for time-series data
- 2018 - 2.0 of InfluxData

Why InfluxData

- Easy to get started with
- Aims to solve the entire time-series problem
- Scales well
  - Both horizontally and vertically

InfluxDB Data Model

Canonical Time-Series Line Graph

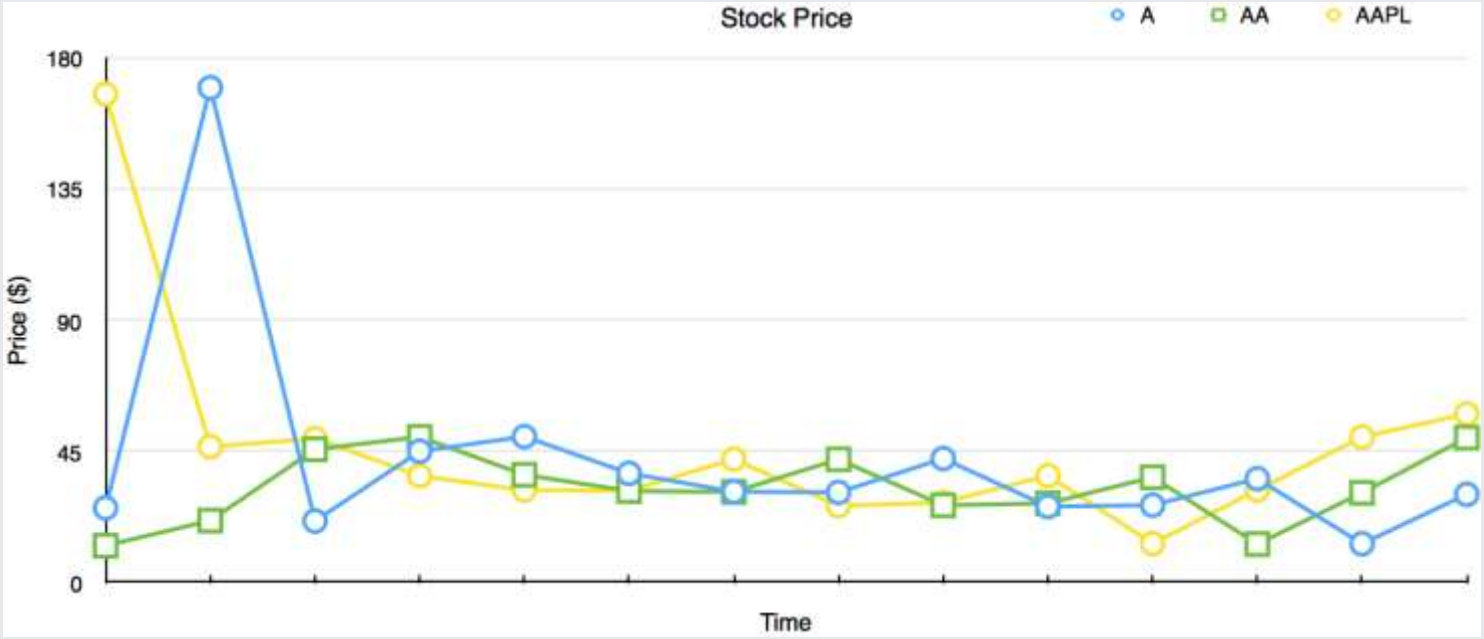

The Label (measurement)

The Legend (tags)

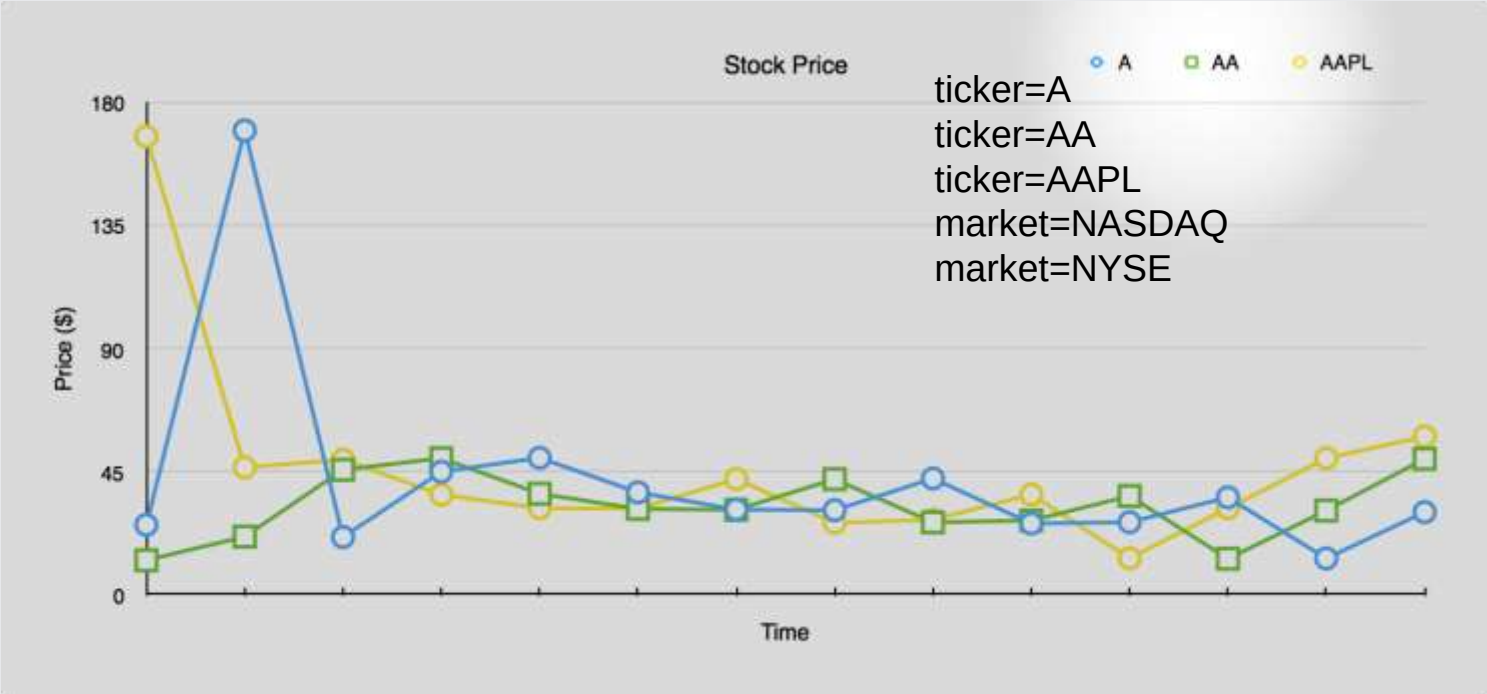

Y-Axis Values

X-Axis Values

Series

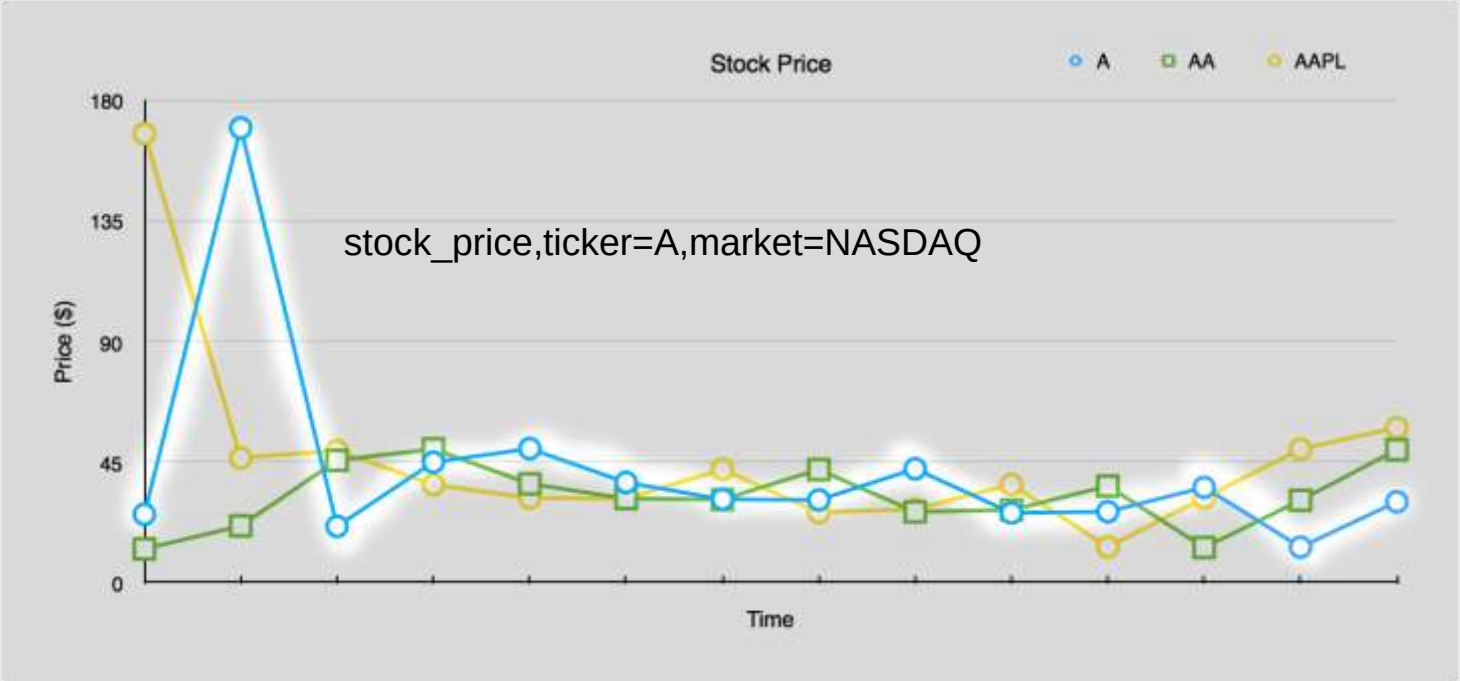
Stock Price
A
AA
AAPL
180
135
90
45
0
Price ($)
Time
stock_price,ticker=A,market=NASDAQ

Data Model

- Measurement
  - High level grouping of data
- Tags
  - Indexed key-value pairs
- Fields
  - Key-value pairs of actual data
- Timestamp
  - Time of the data
- Series
  - A unique combination of measurement+tags

Line Protocol

cpu,host=serverA,num=1,region=west idle=1.667,system=2342.2 1492214400000000000

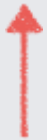

Measurement

Line Protocol

cpu,host=serverA,num=1,region=west idle=1.667,system=2342.2 1492214400000000000

Line Protocol

cpu,host=serverA,num=1,region=west idle=1.667,system=2342.2 1492214400000000000

Fields

Line Protocol

cpu,host=serverA,num=1,region=west idle=1.667,system=2342.2 1492214400000000000

timestamp

Querying Data

InfluxData Languages

- InfluxQL
  - SQL-like query language
- TICKscript
  - Time-series data processing language
- Flux
  - Next generation functional data scripting language

InfluxQL

```
> SELECT index, id FROM h2o_quality WHERE time > now() - 1w GROUP BY location
name: h2o_quality
tags: location = coyote_creek
time                  index  id
----                  -----  ---
2015-08-18T00:00:00Z  41      1
2015-08-18T00:00:00Z  41      1

name: h2o_quality
tags: location = santa_monica
time                  index  id
----                  -----  ---
2015-08-18T00:00:00Z  99     2
2015-08-18T00:06:00Z  56     2
```

TICKscript

```
var measurement = 'requests'

var data = stream
    |from()
        .measurement(measurement)
    |where(lambda: "is_up" == TRUE)
    |where(lambda: "my_field" > 10)
    |window()
         .period(5m)
         .every(5m)

// Count number of points in window
data
     |count('value')
       .as('the_count')

// Compute mean of data window
data
     |mean('value')
      .as('the_average')
```

Flux

```
// get all data from the telegraf db
from(bucket:”telegraf/autogen”)
  // filter that by the last hour
  |> range(start:-1h)
  // filter further by series with a specific measurement and field
  |> filter(fn: (r) => r._measurement == "cpu" and r._field == "usage_system")
```

Why Flux?

- Composabile
  - Users should be able to take pieces of different scripts and combine them into a single one to solve their own problem
- Extensible
  - Adding new functions and capabilities to flux should be easy
- Shareable
  - Users should be able to create libraries and packages to solve specific problems
- Flexible
  - Users should be able to use the language for arbitrary data processing

Schema Design

What not to do

Don’t Encode Data into Measurement/Tags

Bad:

```
cpu.server-5.us-west value=2 1444234982000000000

cpu.server-6.us-west value=4 1444234982000000000

mem-free.server-6.us-west value=2500 1444234982000000000
```

Good:

```
cpu,host=server-5,region=us-west value=2 1444234982000000000

cpu,host=server-6,region=us-west value=4 1444234982000000000

mem-free,host=server-6,region=us-west value=2500 1444234982000000
```

Don’t Encode Data into Measurement/Tags

Bad:

```
cpu,server=localhost.us-west value=2 1444234982000000000

cpu,server=localhost.us-east value=3 1444234982000000000
```

Good:

```
cpu,host=localhost,region=us-west value=2 1444234982000000000

cpu,host=localhost,region=us-east value=3 1444234982000000000
```

Don’t Use Tags with Very High Variability

Bad:

```
response_time,session_id=33254331,request_id=3424347 value=340 14442349820000
```

Good-ish:

```
response_time session_id=33254331,request_id=3424347,value=340 14442349820000
```

Don’t Use Too Few Tags

Bad:

```
cpu,region=us-west host="server1",value=0.5 1444234986000
cpu,region=us-west host="server2",value=4 1444234982000
cpu,region=us-west host="server2",value=1 1444234982000
```

Good-ish:

```
cpu,region=us-west,host=server1 value=0.5 1444234986000
cpu,region=us-west,host=server2 value=4 1444234982000
cpu,region=us-west,host=server2 value=1 1444234982000
```

What should I do then?

Designing a Schema

- What dashboards do I need?
- What alerts do I need?
- What kind of reports do I want to generate?
- What type of information do I need readily available when there's an incident?

Schema Example

- I operate a SaaS application
- There are ~1000 different services
- I want to know the request and error rates for each service
- I want to trigger an alert if the error rate for each service
- I want to see the services with the highest average request duration

Data Available

- `app` Service name, e.g. user_service, auth_service…
- `container_id` Container ID of the container running the service
- `path` HTTP request path
- `method` HTTP method, e.g. GET, POST, DELETE…
- `src` Hostname of client making request
- `dest` Hostname of server being contacted
- `status` HTTP status code associated with the request
- `request_id` Unique request identifier
- `duration` Duration of request
- `bytes_tx` Number of bytes transmitted
- `bytes_rx` Number of bytes received

Question

Why would it be a bad idea to make `container_id` or `request_id` a tag?

Answer

Why would it be a bad idea to make `container_id` or `request_id` a tag?

`request_id` and `container_id` both have a high cardinality and could result in an large number of series, which impacts memory utilization.

`request_id` is substantially worse than `container_id`. In the next few releases we hope to allow for indexing on `container_id`.

Question

How should we organize our data?

Data Available

- `app` Service name, e.g. user_service, auth_service…
- `container_id` Container ID of the container running the service
- `path` HTTP request path
- `method` HTTP method, e.g. GET, POST, DELETE…
- `src` Hostname of client making request
- `dest` Hostname of server being contacted
- `status` HTTP status code associated with the request
- `request_id` Unique request identifier
- `duration` Duration of request
- `bytes_tx` Number of bytes transmitted
- `bytes_rx` Number of bytes received

Schema

measurement:

```
latency
```

tags:

```
app container_id  path  method  src  dst   status
```

fields:

```
request_id  duration  bytes_tx   bytes_rx
```

Request/Error Rate per Service

Top 10 average request duration

```
> SELECT top(avg_duration, app, 10) FROM (
       SELECT mean(duration) AS avg_dur
              FROM latency
              WHERE time > now() - 1h
              GROUP BY time(1m), *
  )
```

Request Rate Per Service

```
> SELECT count(duration)
  FROM latency
  WHERE time > now() - 10m
  GROUP BY app, time(1s) fill(none)
```

Error Rate Per Service

```
> SELECT count(duration)
  FROM latency
  WHERE time > now() - 10m AND status != '200'
  GROUP BY app, time(1s) fill(none)
```


Thank You!
influx days


2015-06-11T20:46:02Z
1985-10-25T19:28:07.580664347Z
influx days