

Design of Topics and Partitions

Topic Design

- Name
- Schema
- Payload (Data)
- Key
- Number of partitions
- Number of replicas

The Short Story

- DevOps concerns
 - Bandwidth consumption → Size of messages, `serde`, etc.
 - Fault tolerance and availability → Size of cluster, replication factor, etc.
 - Performance → Partitions, message size, cost of serialization, etc.
- Producer concerns
 - Ease of production → Clear schema, cost of serialization, delivery guarantees, etc.
- Consumer concerns
 - Can I subscribe to only what I need → Topics and partitions
 - Latency → Cluster size, performance, etc.

How Topics and Partitions Influence Concerns?

- Topic topology
 - Schema (structure, format, etc.)
 - Temporal constraints (frequency, triggers, etc.)
 - Do you use topics to allow for fine grain subscriptions?
- Partitions
 - Determines throughput (but not without cost)
 - Can be used for fine-grained subscription (requires use of low level API)
- Recommendation
 - Use topics to convey semantic
 - Use partition to control throughput

Name

- Descriptive name
- Don't hardcode the name all over your application!
- Rule of thumb:
 - Use a longer name that is easy to understand

Schema

- JSON

- Common choice
- Not the most efficient

- Apache Avro

- Binary format
- Compression
- Schema evolution
- Dynamic typing (no code generation needed for serialization)



Partitions and Throughput

- Unit of parallelism: topic partition
- Writes to different partitions done in parallel
- Consumer: one thread get a single partition's data
- Consumer parallelism: bounded by the number of consumed partitions
- Throughput on a producer is a function of:
 - Batching size
 - Compression codec
 - Acknowledgement type
 - Replication factor
- Consumer throughput is a function of the message processing logic

Overpartitioning

- Problem: Increasing the number of partition and message ordering
 - If messages have keys, increasing the number of partitions may cause problems
 - Kafka maps a message to a partition based on the hash of the key
 - Messages with the same key go to the same partition
 - If we increase the number of partitions, this does not hold
 - Messages with the same key may for the retention period appear in multiple partitions
- Therefore:
 - Overpartition for a situation you expect in future

Too Many Partitions

- Each partition maps to a directory in the broker
 - 2 files: index, actual data
- You may need to configure the open file handle limit
 - Configuration
 - Seen in production > 30K open file handles / broker

Partitions and Availability

- Intra-cluster replication
- A partition can have multiple replicas, each on a different broker
- Clean broker shutdown: the controller moves the leaders off the broker that is shutting down
 - Takes a couple of ms
- Unclean shutdown: the loss of availability is dependent on the number of partitions
 - All replicas become unavailable at the same time
 - A new leader must be elected
 - Potential unavailability in seconds

Partitions and End-to-End Latency

- Consumers can consume a message after it has been committed
 - Requires replication to all in-sync replicas
- Commit time can be significant
 - Too long for some real-time systems
- A rule of thumb:
 - Limit the number of partitions per broker to:
 $100 \times (\text{number of brokers}) \times (\text{replication factor})$

Based on [Confluent Blog](#)

Partitions and Client Memory

- A producer can set the amount of memory for buffering messages
 - Messages are buffered per partition
 - When buffer is full, messages are sent to the broker
- More partitions: more message buffering in the producer
- If out of memory, producer will block or drop new messages
 - Reconfigure
- Allocate at least a few tens of KB per partition

Summary

- Design topics based on your application semantics
 - Message types
 - Consumer concerns
 - Subscription granularity
- Decide on the number of partitions based on throughput
 - The more partitions, the higher theoretical throughput
 - Not without penalty
 - File handlers, consumer memory, latency, etc
 - Evaluate to overpartition to accommodate future growth